

Multiprocessor Real-Time Scheduling

Embedded System Software Design

Prof. Ya-Shu Chen
National Taiwan University of
Science and Technology

Outline

- Multiprocessor Real-Time Scheduling
- Global Scheduling
- Partitioned Scheduling
- Semi-partitioned Scheduling

Multiprocessor Models

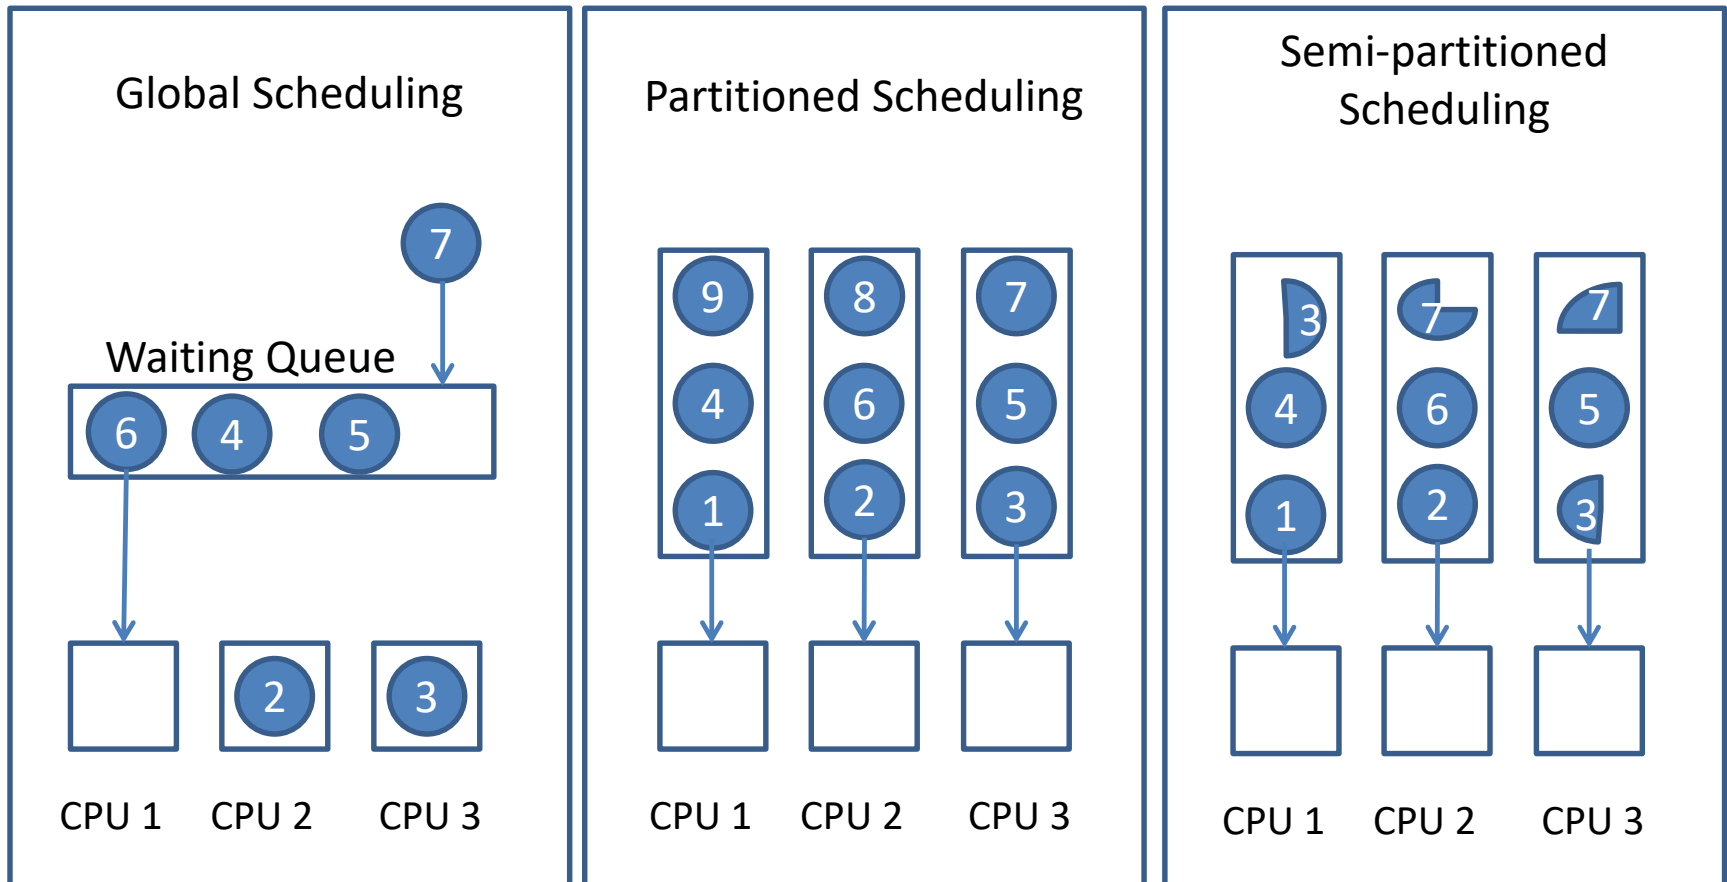
先考慮每一個核心性能一樣

- **Identical (Homogeneous):** All the processors have the same characteristics, i.e., the execution time of a job is independent on the processor it is executed.
- **Uniform:** Each processor has its own speed, i.e., the execution time of a job on a processor is proportional to the speed of the processor.
 - A faster processor always executes a job faster than slow processors do.
 - For example, multiprocessors with the same instruction set but with different supply voltages/frequencies.
- **Unrelated (Heterogeneous):** Each job has its own execution time on a specified processor
 - A job might be executed faster on a processor, but other jobs might be slower on that processor.
 - For example, multiprocessors with different instruction sets.

Scheduling Models

- Global Scheduling:
 - A job may execute on any processor.
 - The system maintains a global ready queue.
 - Execute the M highest-priority jobs in the ready queue, where M is the number of processors.
 - It requires high on-line overhead.
- Partitioned Scheduling:
 - Each task is assigned on a dedicated processor.
 - Schedulability is done individually on each processor.
 - It requires no additional on-line overhead.
- Semi-partitioned Scheduling:
 - Adopt task partitioning first and reserve time slots (bandwidths) for tasks that allow migration.
 - It requires some on-line overhead.

Scheduling Models



Global Scheduling

- All ready tasks are kept in a global queue
- A job can be migrated to any processor. task 可以在任何一個core上執行
- Priority-based global scheduling:
 - Among the jobs in the global queue, the M highest priority jobs are chosen to be executed on M processors.
 - Task migration here is assumed with no overhead.
- Global-EDF: When a job finishes or arrives to the global queue, the M jobs in the queue with the shortest absolute deadlines are chosen to be executed on M processors. 一樣先用EDF排出priorid
- Global-RM: When a job finishes or arrives to the global queue, the M jobs in the queue with the highest priorities are chosen to be executed on M processors.

先用RM 決定 priorid 在分配到有空的cpu

有人做完就找下一個高priorid的人

Global Scheduling

- Advantages:
 - Effective utilization of processing resources (if it works)
 - Unused processor time can easily be reclaimed at run-time (mixture of hard and soft RT tasks to optimize resource utilization)
讓資源有效的使用, 我哪裡都可以去
可以比較容易 可排成,
hard 不能miss, soft可以但不miss對系統更好
- Disadvantages:
 - Adding processors and reducing computation times and other parameters can actually decrease optimal performance in some scenarios!
把CPU資源變多
在有效時間內會發現
其實資源使用效率沒有比較好.
 - Poor resource utilization for hard timing constraints
 - Few results from single-processor scheduling can be used
資源使用會受時間限制
global 沒辦法用沿用single processor 結果
partition 不會被洗 因為都在同一個core , 所以用single processor就可以

Schedule Anomaly

異常

- **Anomaly 1**


A decrease in processor demand from higher-priority tasks can *increase* the interference on a lower-priority task because of the change in the time when the tasks execute

- **Anomaly 2**

A decrease in processor demand of a task *negatively* affects the task itself because the change in the task arrival times make it suffer more interference


Anomaly 1

task set: schedulable



task set: unschedulable

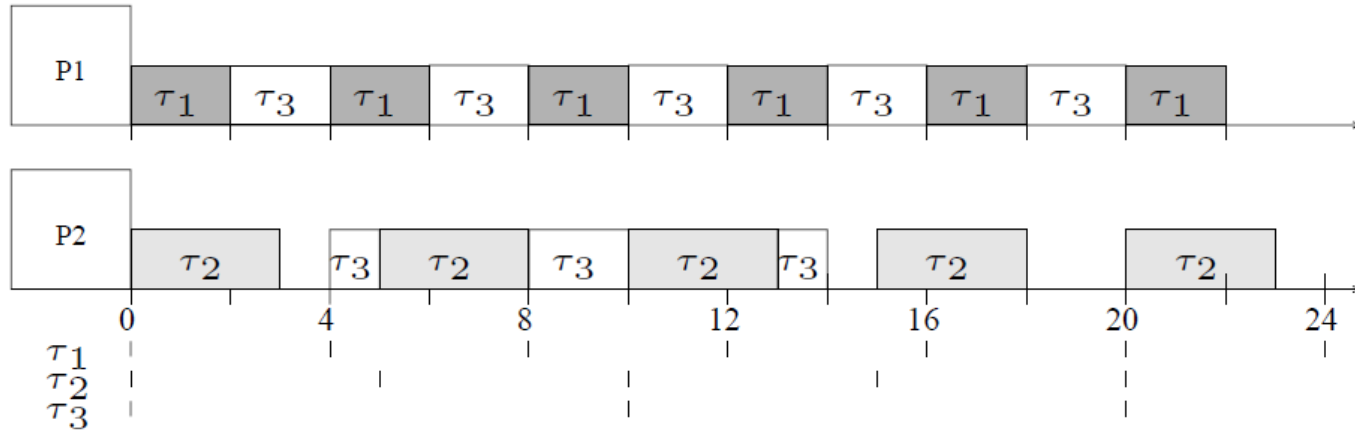
週期拉長了 反而T3不能做 讓他做不完了




τ_3 needs to execute two more time units.

Anomaly 2

task set: schedulable



task set: unschedulable



使用率變小了 但排不了 因為來的時間不對
自己也會影響自己了

τ_3 needs to execute
one more time unit.

Dhall effect

- Dhall effect : For Global-EDF or Global-RM, the least upper bound for schedulability analysis is at most 1.
- On 2 processors:

Task	T	D	C	U
T1	10	10	5	0.5
T2	10	10	5	0.5
T3	12	12	8	0.67

- T3 is not schedulable

先partition
在schedulable

有兩顆CPU
但可能只有一顆可以用

Schedulability Test

- A set of periodic tasks t_1, t_2, \dots, t_N with implicit deadlines is schedulable on M processors by using preemptive Global EDF scheduling if

最大的U會影響整個可不可以執行

$$\sum_{i=1}^N \frac{C_i}{T_i} \leq M \left(1 - \frac{C_k}{T_k}\right) + \frac{C_k}{T_k},$$

where t_k is the task with the largest utilization C_k/T_k

M幾個processor
以及跟最大的U有關

Weakness of Global Scheduling

- Migration overhead 重新搬移
- Schedule Anomaly

Partitioned Scheduling

分群
一顆processors Us就是一

- Two steps:
 - Determine a mapping of tasks to processors
 - Perform run-time single-processor scheduling
- Partitioned with EDF
 - Assign tasks to the processors such that no processor's capacity is exceeded (utilization bounded by 1.0)
 - Schedule each processor using EDF

決定Map到哪個processor
決定是在哪個cpu上做

用single processor去排成

Bin-packing Problem

Given a bin size V and a list a_1, \dots, a_n of sizes of the items to pack, find an integer B and a B -partition $S_1 \cup \dots \cup S_B$ of $\{1, \dots, n\}$ such that $\sum_{i \in S_k} a_i \leq V$, for all $k = 1, \dots, B$.
A solution is optimal if it has minimal B .

The problem is NP-complete !!

Bin-packing to Multiprocessor Scheduling

- The problem concerns packing objects of varying sizes in boxes ("bins") with the objective of minimizing number of used boxes.
 - Solutions (Heuristics): First Fit
- Application to multiprocessor systems:
 - Bins are represented by processors and objects by tasks.
 - The decision whether a processor is "full" or not is derived from a utilization-based schedulability test.

Partitioned Scheduling

- Advantages:
 - Most techniques for single-processor scheduling are also applicable here
- Partitioning of tasks can be automated
 - Solving a bin-packing algorithm
- Disadvantages:
 - Cannot exploit/share all unused processor time
 - May have very low utilization, bounded by 50%

Partitioned Scheduling Problem

Given a set of tasks with arbitrary deadlines, the objective is to decide a feasible task assignment onto M processors such that all the tasks meet their timing constraints, where C_i is the execution time of task t_i on any processor m .

Partitioned Algorithm

- First-Fit: choose the one with the smallest index
找最小 processor 開始放
- Best-Fit: choose the one with the maximal utilization
找最大U來放
- Worst-Fit: choose the one with the minimal utilization
找最小U來放

Partitioned Example

Utilization

- 0.2 -> 0.6 -> 0.4 -> 0.7 -> 0.1 -> 0.3




First Fit

比較直覺的作法



Best Fit

盡可能不公平
用到的core可以更少



Worst Fit

盡量平衡
整個系統的平行度比較好

EDF with First Fit

Input: A task set $\tau = \{\tau_1, \tau_2, \dots, \tau_n\}$ and a set of processors $\{p_1, \dots, p_m\}$

Output: j ; number of processors required.

```
1.   $i := 1; j := 1; k_q = 0; (\forall_q)$ 
2.  while ( $i \leq n$ ) do
3.       $q := 1;$ 
4.      while ( $(U_q + u_i) > 1$ ) do    edf 改rm 要改成rm 的U
5.           $q := q + 1; /* \text{increase the processor index} */$ 
6.           $U_q := U_q + u_i; k_q := k_q + 1;$ 
7.          if ( $q > j$ ) then
8.               $j := q;$ 
9.           $i := i + 1;$ 
10. return ( $j$ );     $j$  需要幾個processor
11. end
```

Schedulability Test

Lopez [3] proves that the **worst-case achievable utilization** for EDF scheduling and FF allocation (EDF-FF) takes the value

0.5 0.5 0.67
2 process

processor 的數量
還有 U

If all the tasks have an utilization factor C/T under a value α , where m is the number of processors

$$U_{wc}^{EDF-FF}(m, \beta) = \frac{\beta m + 1}{\beta + 1} \quad \text{where} \quad \beta = \lfloor 1/\alpha \rfloor$$

最大的utilization

Demand Bound Function

- Define demand bound function $dbf(\tau_i, t)$ as

$$dbf(\tau_i, t) = \max \left\{ 0, \left\lfloor \frac{t + T_i - D_i}{T_i} \right\rfloor \right\} C_i = \max \left\{ 0, \left\lfloor \frac{t - D_i}{T_i} \right\rfloor + 1 \right\} C_i.$$


T_i 最小間隔

- We need approximation to enforce polynomial-time schedulability test

$$dbf^*(\tau_i, t) = \begin{cases} 0 & \text{if } t < D_i \\ (\frac{t-D_i}{T_i} + 1) C_i & \text{otherwise.} \end{cases}$$

deadline

執行時間



Deadline Monotonic Partition

Input: \mathbf{T}, M ;

- 1: re-index (sort) tasks such that $D_i \leq D_j$ for $i < j$;
- 2: $\mathbf{T}_i \leftarrow \emptyset, U_i \leftarrow 0, \forall m = 1, 2, \dots, M$;
- 3: **for** $i = 1$ to N , where $N = |\mathbf{T}|$ **do** 幾個task
- 4: **for** $m = 1$ to M **do** 幾個core 因為他是非週期要加DBF
- 5: **if** $\frac{C_i}{T_i} + \sum_{\tau_j \in \mathbf{T}_m} \frac{C_j}{T_j} \leq 1$ and $C_i + \sum_{\tau_j \in \mathbf{T}_m} dbf^*(\tau_j, D_i) \leq D_i$ **then** 因為deadline period可能不相同
- 6: assign task τ_i onto processor m and $\mathbf{T}_m \leftarrow \mathbf{T}_m \cup \{\tau_i\}$; 看看可不可以排進去
- 7: break;
- 8: **if** τ_i is not assigned **then**
- 9: return "The task assignment fails"; 如果都排不進去
- 10: return feasible task assignment $\mathbf{T}_1, \mathbf{T}_2, \dots, \mathbf{T}_M$;

Schedulability Test

Theorem 4 *Any sporadic task system τ is successfully scheduled by Algorithm PARTITION on m unit-capacity processors, for any*

$$m \geq \left(\frac{2\delta_{\text{sum}} - \delta_{\text{max}}}{1 - \delta_{\text{max}}} + \frac{u_{\text{sum}} - u_{\text{max}}}{1 - u_{\text{max}}} \right) \quad (14)$$

$$\begin{array}{ll} \delta_{\text{max}} & \stackrel{\text{def}}{=} \max_{i=1}^n (e_i/d_i) & u_{\text{max}} & \stackrel{\text{def}}{=} \max_{i=1}^n (u_i) \\ \delta_{\text{sum}} & \stackrel{\text{def}}{=} \max_{t>0} \left(\frac{\sum_{j=1}^n \text{DBF}(\tau_j, t)}{t} \right) & u_{\text{sum}} & \stackrel{\text{def}}{=} \sum_{j=1}^n u_j \end{array}$$

Weakness of Partitioned Scheduling

- Restricting a task on a processor reduces the schedulability
- Restricting a task on a processor makes the problem NP-hard
- Example: Suppose that there are M processors and $M + 1$ tasks with the same period T and the (worst-case) execution times of all these $M + 1$ tasks are $T/2 + e$ with $e > 0$
 - With partitioned scheduling, it is not schedulable



Semi-partitioned Scheduling

可以切進去就切進去 分割來排
但還是不能平行同時做
只是任務切割

- Tasks are first partitioned into processor.
- To reduce the utilization, we again pick the processor with the minimum task utilization
- If a task cannot fit into the picked processor, we will have to split it into multiple (two or even more) parts.
- If t_i is split and assigned to a processor m and the utilization on processor m after assigning t_i is at most $U(\text{scheduler}, N)$, then t_i is so far schedulable.

Semi-partitioned EDF


- T_{\min} is the minimum period among all the tasks.
- By a user-designed parameter k , we divide time into slots with length $S = T_{\min}/k$. 切成幾個區塊 K
- We can use the first-fit approach by splitting a task into 2 subtasks, in which one is executed on processor m and the other is executed on processor $m + 1$.
- Execution of a split task is only possible in the reserved time window in the time slot.
- Applying first-fit algorithm, by taking SEP as the upper bound of utilization on a processor.
- If a task does not fit, split this task into two subtasks and allocate a new processor, one is assigned on the processor under consideration, and the other is assigned on the newly allocated processor.



一直換來換去
要一直搬資料
context switch 會很多次


Semi-partitioned EDF

- For each time slot, we will reserve two parts.



If a task t_i is split, the task can be served only within these two pre-defined time slots with length x_i and y_i .

很難剛好切一半
後半段要做要等前半段做完



還要考慮比他高優先權的人

A processor can host two split tasks, t_i and t_j . t_i is served at the beginning of the time slot, and t_j is served at the end.

The schedule is EDF, but if a split task instance is in the ready queue, it is executed in the reserved time region.

排成 Semi-partitioned EDF

if sep = 0.9

- We can assign all the tasks t_i with $U_i > SEP$ on a dedicated processor. So, we only consider tasks with U_i no larger SEP.

```
1:  $m \leftarrow 1, U_m \leftarrow 0;$ 
2: for  $i = 1$  to  $N$ , where  $N = |\mathbf{T}|$  do
3:   if  $\frac{C_i}{T_i} + U_m \leq SEP$  then
4:     assign task  $\tau_i$  on processor  $m$ ;
5:      $U_m \leftarrow U_m + \frac{C_i}{T_i};$ 
6:   else
7:     assign task  $\tau_i$  on processor  $m$  with  $lo\_split(\tau_i)$  set to  $SEP - U_m$  and on
       processor  $m + 1$  with  $high\_split(\tau_i)$  set to  $\frac{C_i}{T_i} - (SEP - U_m);$ 
8:      $m \leftarrow m + 1$  and  $U_m \leftarrow \frac{C_i}{T_i} - (SEP - U_m);$ 
```

When executing, the reservation to serve t_i is to set x_i to $S \times (f + lo_split(t_i))$ and y_i to $S \times (f + high_split(t_i))$.
SEP is set as a constant.

Two Split Tasks on a Processor

- For split tasks to be schedulable, the following sufficient conditions have to be satisfied
 - $\text{lo_split}(t_i) + f + \text{high_split}(t_i) + f \leq 1$ for any split task t_i . 被切成兩塊你的F是多少
 - $\text{lo_split}(t_j) + f + \text{high_split}(t_i) + f \leq 1$ when t_i and t_j are assigned on the same processor.
- Therefore, the “magic value” SEP

$$SEP \leq 1 - 2f \leq 1 - 2(\sqrt[2]{\kappa(\kappa + 1)} - \kappa).$$

- However, we still have to guarantee the schedulability of the non-split tasks. It can be shown that the sufficient condition is

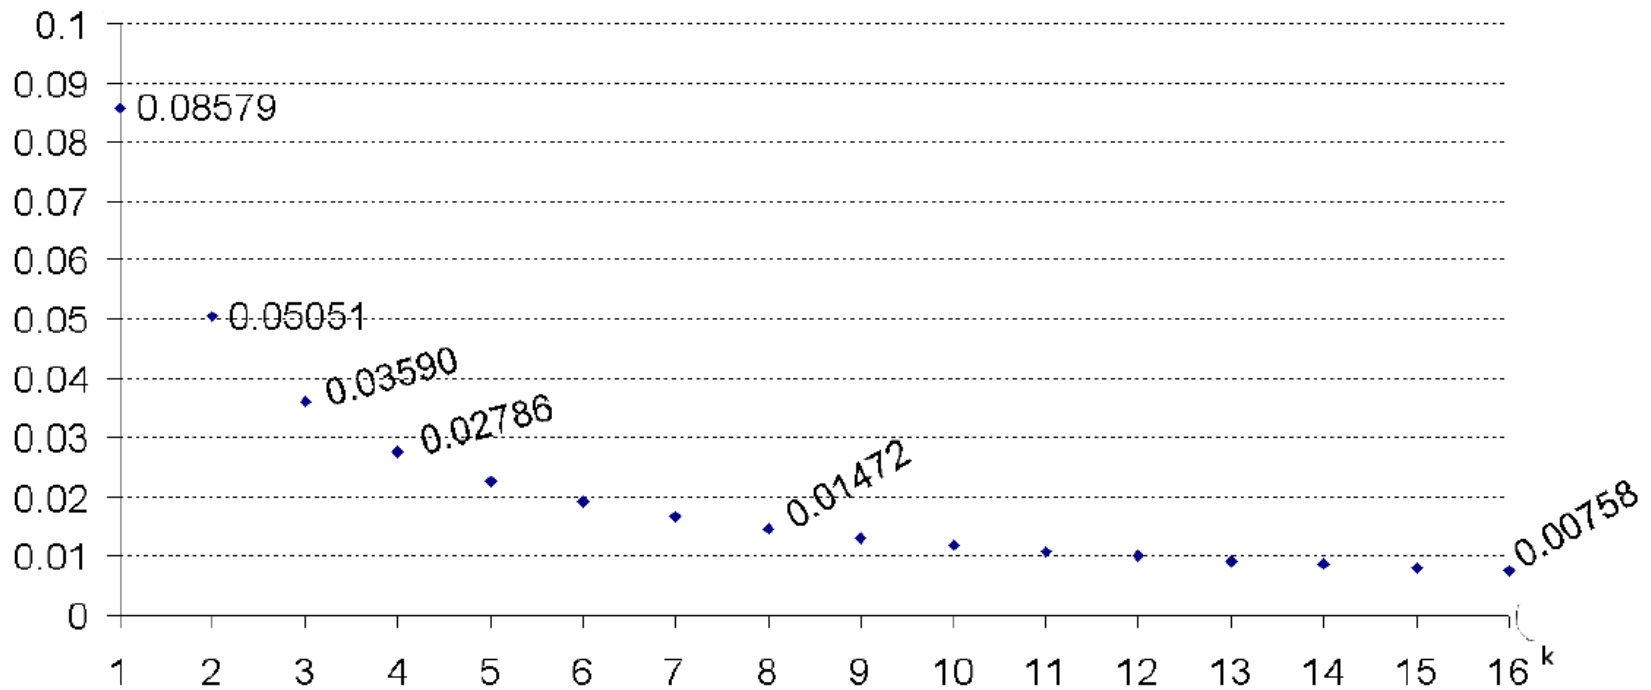
$$SEP \leq 1 - 4f \leq 1 - 4(\sqrt[2]{\kappa(\kappa + 1)} - \kappa).$$

Schedulability Test

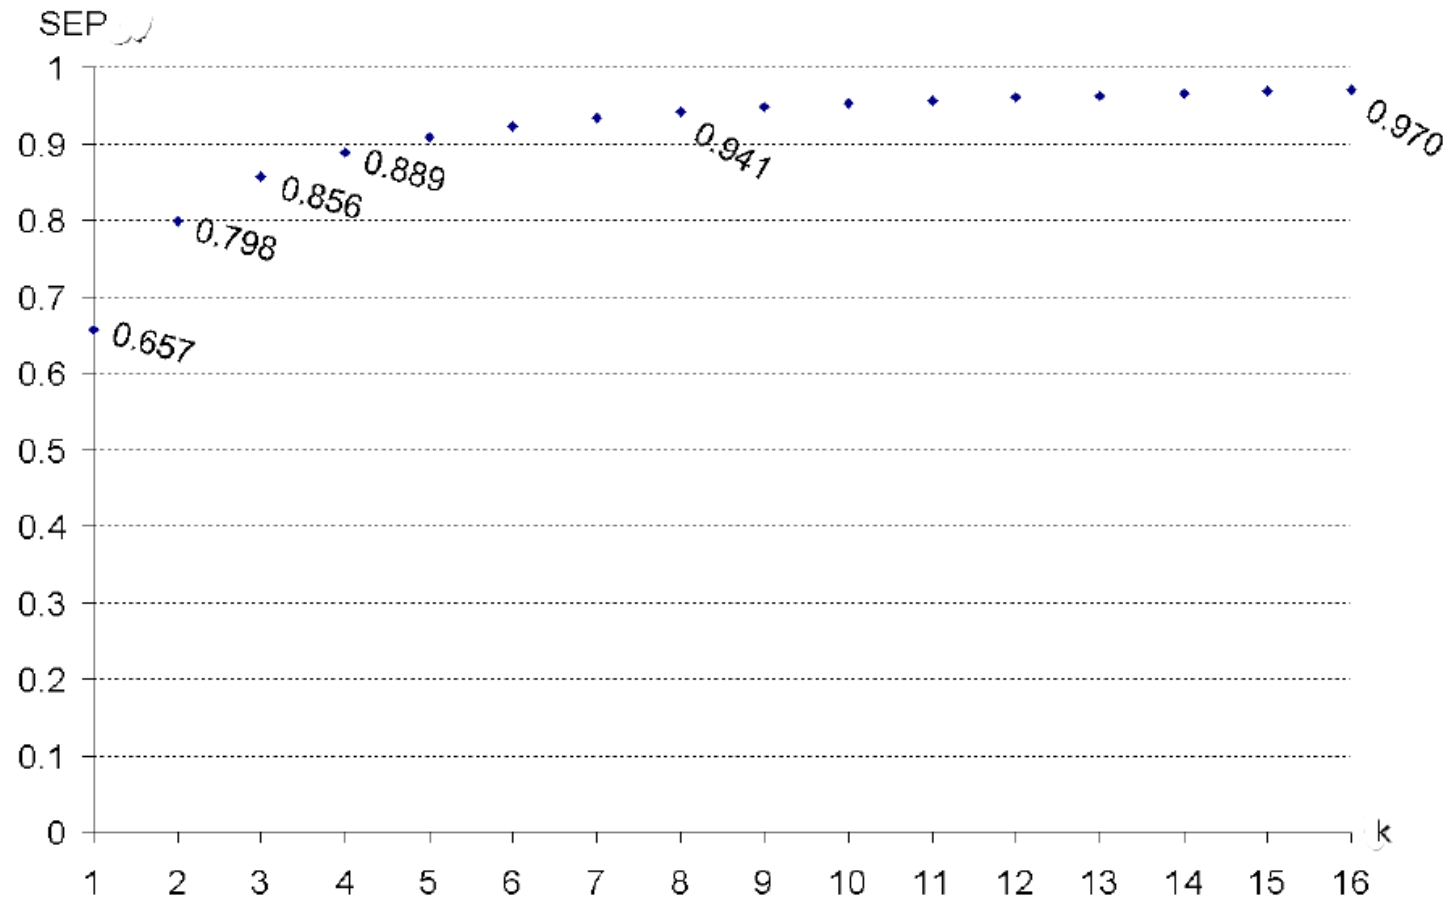
By taking SEP as $1 - 4(\sqrt[2]{\kappa(\kappa + 1)} - \kappa)$ and $f = \sqrt[2]{\kappa(\kappa + 1)} - \kappa$, the above algorithm guarantees to derive feasible schedule if $\sum_{\tau_i \in \mathbf{T}} \frac{C_i}{T_i} \leq M' \cdot SEP$ and $\frac{C_i}{T_i} \leq SEP$ for all tasks τ_i .

K的證明

Magic Values: f



Magic Values: SEP



Reference

- Multiprocessor Real-Time Scheduling
 - Dr. Jian-Jia Chen: Multiprocessor Scheduling. Karlsruhe Institute of Technology (KIT): 2011-2012
- Global Scheduling
 - Sanjoy K. Baruah: Techniques for Multiprocessor Global Schedulability Analysis. RTSS 2007: 119-128
- Partitioned Scheduling
 - Sanjoy K. Baruah, Nathan Fisher: The Partitioned Multiprocessor Scheduling of Sporadic Task Systems. RTSS 2005: 321-329
- Semi-partitioned Scheduling
 - Björn Andersson, Konstantinos Bletsas: Sporadic Multiprocessor Scheduling with Few Preemptions. ECRTS 2008: 243-252

See You Next Week